

Three bets I would put spend behind

Part 2 of 3 — A/B Test Ideation · [variant-1-score-first.html](#) · [variant-2-eight-questions.html](#) · [variant-3-signup-first.html](#) · [contracts: docs/EXPERIMENT-CONTRACTS.md](#)

Decision. Three architecture-level bets, one variable each, each a standalone copy of the control from one template (the build prints the diff: 14, 32 and 16 changed lines): how much result to give before the ask, how much effort to require, when to ask.

Metric. Account CVR = account creations / page visits, intention-to-treat; decided in production on verified accounts per eligible new visitor within seven days; a variant that wins on raw accounts and loses on verified accounts is a loss.

Trade-off. Every call is read on the two-sided 95% interval (Test 1's non-inferiority margin too), sized on the observed base rate after two weeks of baseline; fixed horizon, daily sample-ratio check, step metrics after the primary call.

TEST 1 — RESULT GATING: DISCLOSED GATE (CONTROL) VS SCORE-FIRST · STAGE: RESULT → SIGNUP

Hypothesis. Score-first raises trust, starts and email validity but removes the reason to sign up at the moment of highest interest, so Account CVR falls; worth its spend because in this distrusted category trust could beat curiosity.

Mechanism. The control gates the range and rank behind a free email-only account announced before question one; the variant shows range and rank the moment the quiz ends and sells the account on the review, history and updating rank; the landing lead and disclosure describe each arm's own gate. **KPIs.** Verified-account CVR decides; step metrics: landing → start, completion, result → signup view, signup view → complete (the gate rate). **Guardrails.** Verified-email rate, disposable-domain share, result views after signup. **Rule and risk.** Ship score-first if the 95% interval on its verified-account CVR is not below the control's by more than 10% relative and guardrails hold; until traffic supports that read, ship it only if it wins outright. Decisive risk: a loss could come from the longer result screen (CTA 888 vs 638 px), a declared, measured coupling.

TEST 2 — QUIZ LENGTH: 12 QUESTIONS (CONTROL) VS 8 · STAGE: QUIZ START AND COMPLETION

Hypothesis. Completion rises by more than the gate rate falls, so Account CVR rises; if not, effort is buying credibility rather than costing it. **Mechanism.** Eight items sampled from the same bank across the same difficulty tiers and areas (two per area, all three difficulty tiers); the copy says "8 questions · ~3 min"; everything else is identical. A friction test more than a pure length test: which items appear and the score resolution change too, so the top band means "at most one wrong" in both arms (11/12, 7/8) and the band distribution is a guardrail. **KPIs.** Verified-account CVR decides; step metrics: start rate, completion, per-question drop-off, median completion time, result → signup view as the credibility signal. **Guardrails.** Band-distribution shift, retake rate, verified-email rate. **Rule and risk.** Ship 8 if the 95% interval lies above the control's and result → signup does not fall by more than 5% relative; a null interval keeps 12. The shared creative says "a short test", not a count, so both arms get the same traffic; the risk is crediting length for composition.

TEST 3 — SIGNUP TIMING: AFTER THE LAST QUESTION (CONTROL) VS BEFORE QUESTION 1 · STAGE: ENTRY

Hypothesis. Raw accounts rise (every starter meets the ask before attrition) while verified accounts and the share of accounts that reach a result fall; the guardrail decides. **Mechanism.** The same free account is created before question one; the landing block, CTA label and signup heading describe the timing. At illustrative magnitudes the late arm yields $0.70 \times 0.75 \times 0.40 \approx 21\%$ of visitors, so the early ask must convert more than a fifth of all visitors. Research briefs split (two call a pre-quiz ask a start-killer; one reports operators profitably moving a paid wall first, a timing analogy only), which is why it earns spend. **KPIs.** Account CVR from the landing view, counting accounts made without a quiz start; step metrics: landing → account, account → quiz start, completion, result views per account. **Guardrails.** Verified-email rate; accounts that reach a result ÷ accounts. **Rule.** Ship if the 95% interval on verified-account CVR lies above the control's by more than 5% relative and at least the control's share of accounts reaches a result; run after Test 1, against the winning gate.

NOT CHOSEN, AND WHERE AI HELPED

A question-1-first landing was the strongest candidate: at 375×667 the disclosure block and a matrix item cannot share the first viewport, so that arm would couple entry architecture with disclosure exposure; it is the first follow-up, then disclosure prominence and a 16-item credibility arm. Claude Code re-scored the backlog against six research inputs, an independent build and an audit, and generated and diffed each variant; the tests, their order, the coupled-copy rules and the decision criteria were my calls.

Labels: [EVIDENCE] · [COMPETITOR OBSERVATION] · [INFERENCE] · [HYPOTHESIS]. Rates above are hypothetical illustrations, not benchmarks.