

Three bets I would put spend behind

Variants: [variant-1-score-first.html](#) · [variant-2-eight-questions.html](#) · [variant-3-signup-first.html](#) — each a complete standalone copy of the control generated from one template; the build prints the diff so isolation is verifiable (14, 32 and 16 changed lines). Live: [yosishe.github.io/iqly-flow](#). Full 24-field contracts: [docs/EXPERIMENT-CONTRACTS.md](#).

Primary KPI for all three: Account CVR = sessions with account creation ÷ sessions with a landing view, intention-to-treat; in production the decision is made on **verified** accounts per eligible new visitor within a seven-day window, with the brief's session metric reported alongside. A variant that wins on raw accounts and loses on verified accounts is a loss. Sizing rule rather than a number: after two weeks of baseline traffic, size each test on the observed base rate for the smallest lift worth shipping (10% relative, 80% power, two-sided, $\alpha = 0.05$) and, where the call is non-inferiority, for a 10% relative margin ([docs/MEASUREMENT-SPEC.md](#)), taking the larger of the two samples; fix the horizon in advance, check sample-ratio mismatch daily, decide on the 95% confidence interval for the primary metric rather than the point estimate, and read step metrics only after the primary call.

Test 1 — Result gating: disclosed gate (control) vs score-first · stage: result → signup (declared effect on landing → start)

Why. The control gates the one thing a "free iq test" searcher came for, the range and rank, behind a free email-only account announced before question one; that is the bet the control depends on most. The variant shows the range, curve and provisional rank free the moment the quiz ends and offers the account for the answer review, the updating rank and saved history; its landing truthfully promises "your IQ range and rank the moment you finish". **Hypothesis.** Score-first raises trust, starts and email validity but removes the reason to sign up at the moment of highest interest. With a hypothetical gate rate of 40% and a "save my result" rate of 12%, starts times completion would have to rise roughly threefold to tie, and no evidence supports a swing that size; it must still be measured, because a category that has taught users to expect a trap could make trust worth more than curiosity. **KPIs.** Account CVR decides, on verified accounts. Step metrics: landing→start, completion, result→signup view, signup view→complete (the gate rate), signup-error rate, and how often the collapsed sections are opened before deciding. **Guardrails and rule.** Verified-email rate, disposable-domain share, result views after signup, complaints. Ship score-first if the 95% interval on its verified-account CVR is not below the control's by more than 10% relative and guardrails hold; until traffic supports that read, ship it only if it wins; keep the gate otherwise. Declared coupling: the landing lead and disclosure must describe each arm's own gate.

Test 2 — Quiz length: 12 questions (control) vs 8 · stage: quiz start and completion

Why. Length acts on the two largest early terms, starts (a shorter honest promise) and completion (fewer break-offs), while acting against the last: a shorter test feels less worth an email, and 8 items are less internally consistent (Cronbach's $\alpha \approx .55$ vs $\approx .65$ for 12, inferred from published short forms). The variant uses eight items sampled across the same difficulty tiers and areas (three easy, two medium, three hard; two per area), so the difficulty ceiling and the area balance are held constant; because the copy must stay true it says "8 questions · ~3 min"; gate, disclosure, signup and result are identical. Primarily a friction test, not a pure length test: length also changes which items appear and the score resolution, so the band distribution is a guardrail, and the top band means at most one wrong in both arms (11/12, 7/8). **Hypothesis.** Completion rises by more than the gate rate falls, so Account CVR rises, because most break-off in a phone test happens in the last, hardest minutes; if it does not, effort is buying credibility rather than costing it. **KPIs.** Account CVR decides. Step metrics: start rate, completion, the per-question drop-off curve, median completion time, result→signup view as the credibility signal. **Guardrails and rule.** Range distribution shift (8/8 is easier than 12/12 with the same cut-points), retake rate, verified-email rate. Ship 8 if the 95% interval on verified-account CVR lies above the control's and result→signup does not fall by more than 5% relative; keep 12 otherwise. In a live run the shared creative would say "a short test" rather than a count, so that the arms receive the same traffic.

Test 3 — Signup timing: after the last question (control) vs before question 1 · stage: entry (landing → account); completion as a quality guardrail

Why. The brief's own example variant (signup after question 3) marks placement as a lever the team considers high-leverage; this is the strongest version of the early-capture bet, chosen over a landing-disclosure dial because it can move the KPI by a multiple rather than a few percent. Every visitor who clicks Start meets the account step before any attrition; the landing block says so ("Free account first"), and the full result appears directly after the last question. An account can exist without a finished quiz, which is exactly why the rule below discounts accounts that never reach a result. **Hypothesis.** Accounts per visit rise if the early ask converts more of all visitors than the late ask converts of finishers: at hypothetical magnitudes the late arm yields $0.70 \times 0.75 \times 0.40 \approx 21\%$ of visitors, so the early ask must convert more than 21% of visitors, roughly a third of those who would have started. Outside opinion splits: two research briefs call an ask before question one a start-killer; one reports operators profitably moving a paid wall first (a timing analogy only); that split is exactly when a test earns spend. Test 1 asks how much result to give before the ask; this test asks when to ask: two levers of the same exchange, run in that order. My prior is that raw accounts rise while verified accounts and the share of accounts that ever reach a result fall. **KPIs.** Account CVR from the landing view, counting the account even when no quiz starts. Step metrics: landing→account (the whole lever), account→quiz start, completion, result views per account. **Guardrails and rule.** Verified-email rate, accounts that reach a result ÷ accounts (an account that never sees a result is not growth), complaints. Ship signup-first if the 95% interval on verified-account CVR lies above the control's by more than 5% relative and at least the control's share of accounts reaches a result; keep the control otherwise. Declared coupling: the landing block, CTA label and signup heading describe the timing; nothing else changes (16 lines).

Not chosen and why. A question-1-first landing was the strongest candidate: measured at 375×667, the disclosure block and a matrix item cannot share the first viewport, so that arm would change entry architecture and disclosure exposure at once (two hypotheses); it is the first follow-up. Then disclosure prominence (block vs one line, both shippable), signup framing (copy-level), and 16 questions as the credibility arm if Test 2 shows the gate rate falling; timer, copy and colour tests cannot move the KPI enough to justify spend. **Where AI helped:** Claude Code re-scored the backlog against two GPT reports, four research briefs, an independent build (Codex) and an audit that argued for a Q1-first third test, generated each variant from the template, diffed the files and ran simulated reviewer passes; the three tests, their order, the coupled-copy rules and the decision criteria were my calls.

Labels: [EVIDENCE] · [COMPETITOR OBSERVATION] · [INFERENCE] · [HYPOTHESIS]. The rates above are hypothetical illustrations of the break-evens, not benchmarks.